Frozen world-model follow-ups | exploratory

Original appearance · horizon 8 · 128 generated roots / 128 recorded episodes

A Physical block-position error

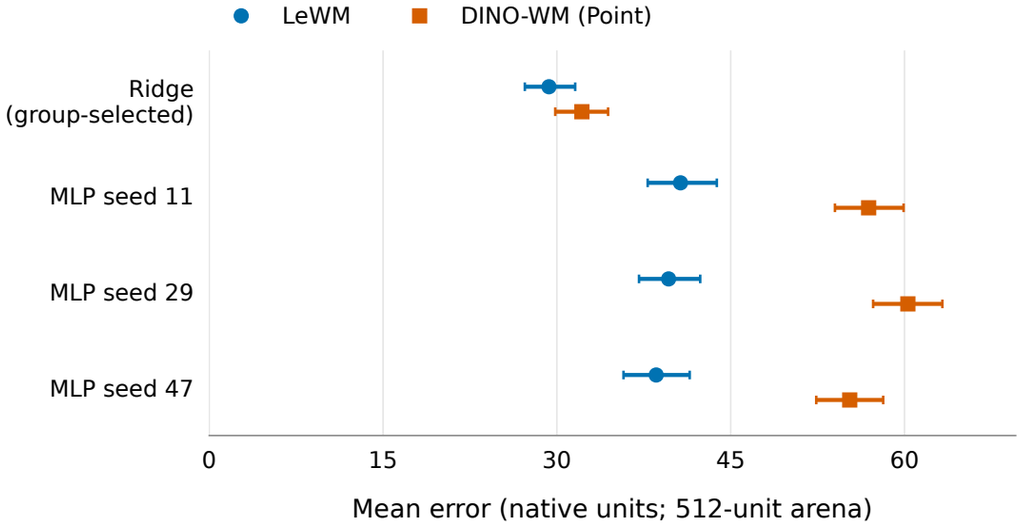

B Actual history − LOCAL persistence

Own axes per model/domain; negative favors actual history

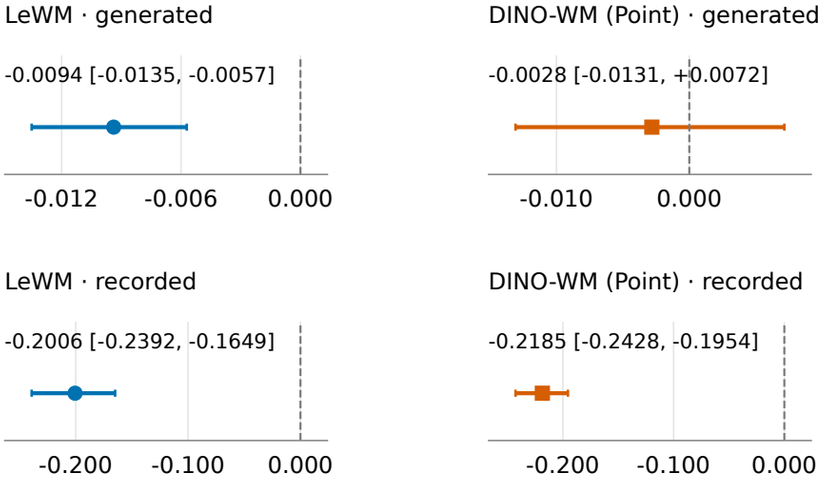

C Action-effect error − zero-effect baseline

Own axes per model; negative favors predicted action effects

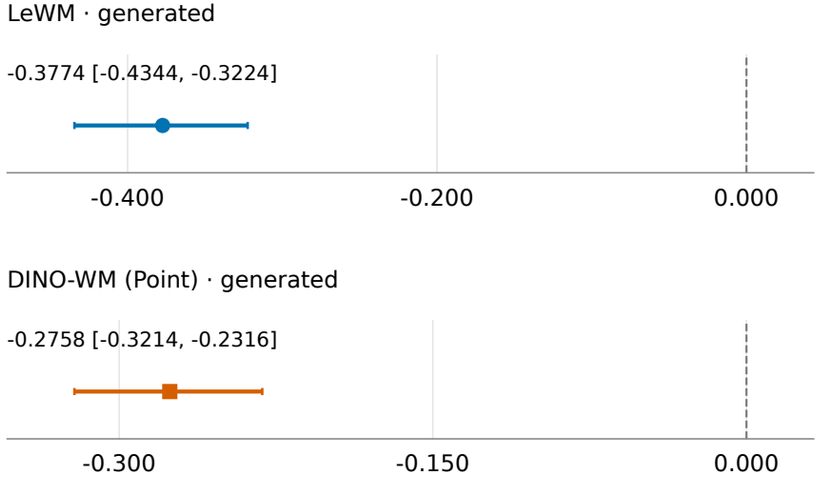

D Finite-bank candidate regret

Mean bank best coverage: 11.28% · 8 candidates

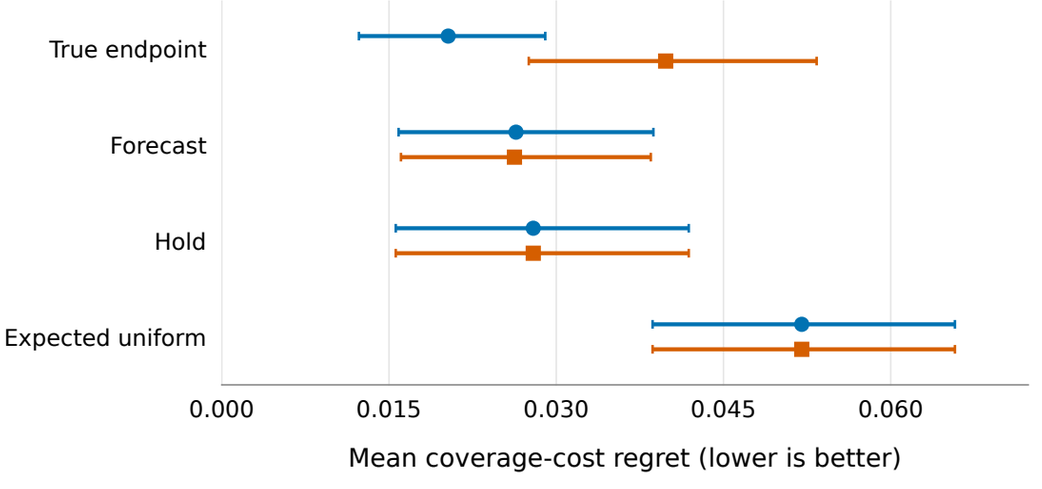

Points and whiskers: pointwise 95% bootstrap intervals (5,000 resamples), equal-root / episode weighting. All three MLP seeds retained.

B/C: checkpoint-specific latent MSE differences; raw scales cannot rank models. Actual history is privileged, not deployable prediction.

Recorded data are pretraining-distribution diagnostics, not held-out generalization. Finite-bank regret is not closed-loop control success.